

Stateless Kafka Brokers


Overview of **architectures** available in this space



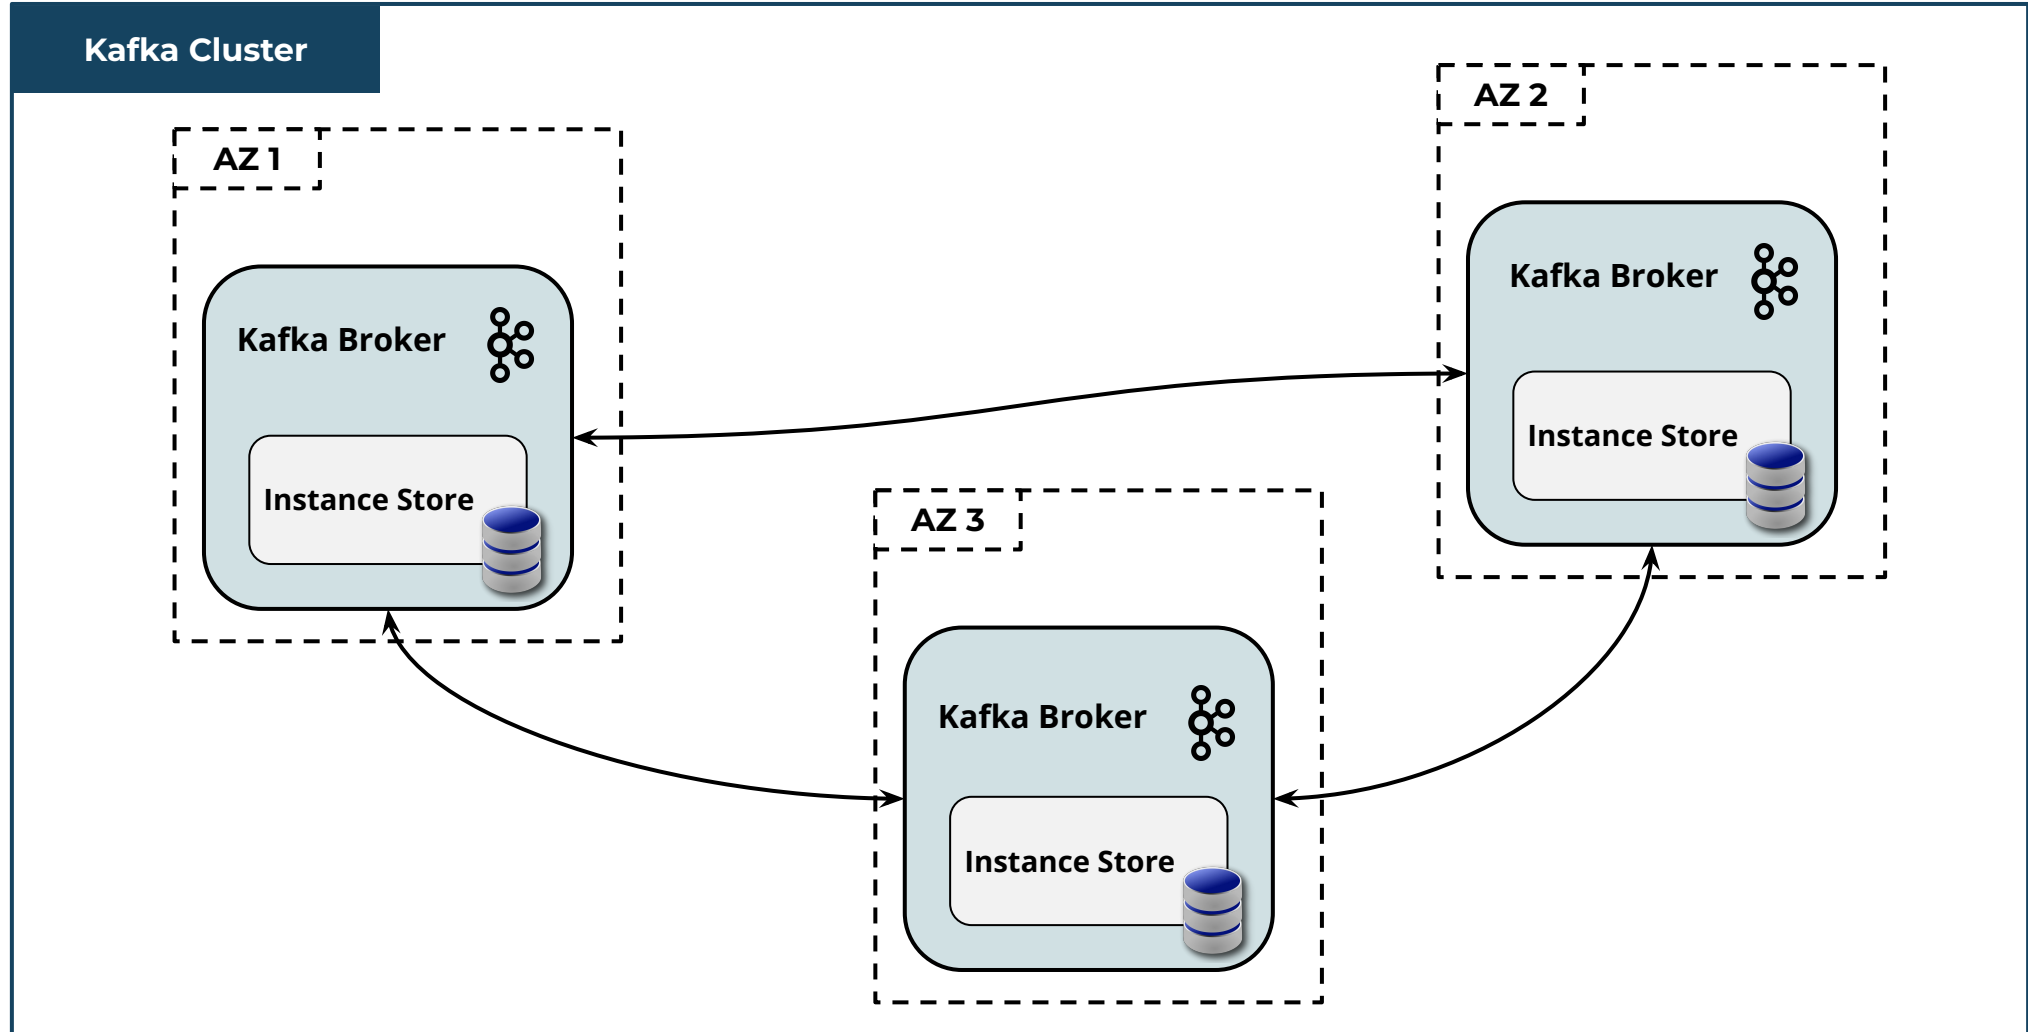
2023




2025



Kafka Highly Available cluster setup



Lack of elasticity means overprovisioning



Source: Confluent

<https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-1-infrastructure/>

Chapter 1:

Apache Kafka with Tiered Storage



How does Tiered Storage work?

block storage

1

2

3

4

5

6

7

8

9

10

11

12



How does Tiered Storage work?

block storage

11

12

object storage

1

2

3

4

5

6

7

8

9

10

11



What's “wrong” with kafka?



Apache Kafka
(w/o Tiered Storage)

- **Cross-az networking**
- **Expensive long-term storage**
- **Poor cluster elasticity**
- **High Operating complexity**



Apache Kafka
+ Tiered Storage

- **Cross-az networking**
- + **Cheap S3 storage**
- + **Elastic scaling**
- + **Simpler to operate**

Where are Kafka's costs?


Cool article by **Confluent** on infrastructure costs of a relatable benchmark

- 100 MBps writes
- 300 MBps reads
- 7 day retention
- **fetch-from-follower not enabled**

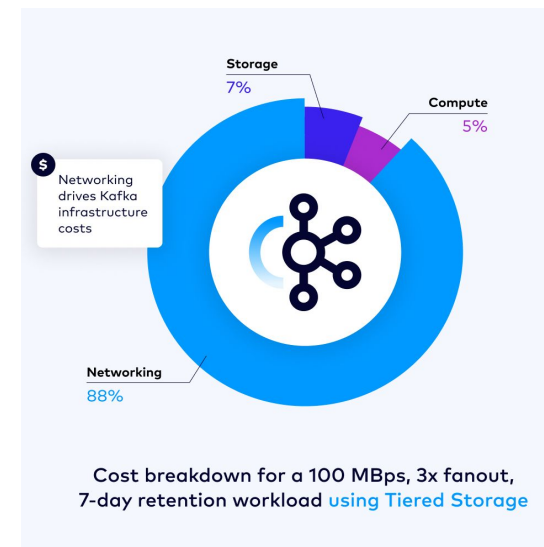
These are **direct costs only**
of a straightforward benchmark

Source: Confluent

<https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-1-infrastructure/>

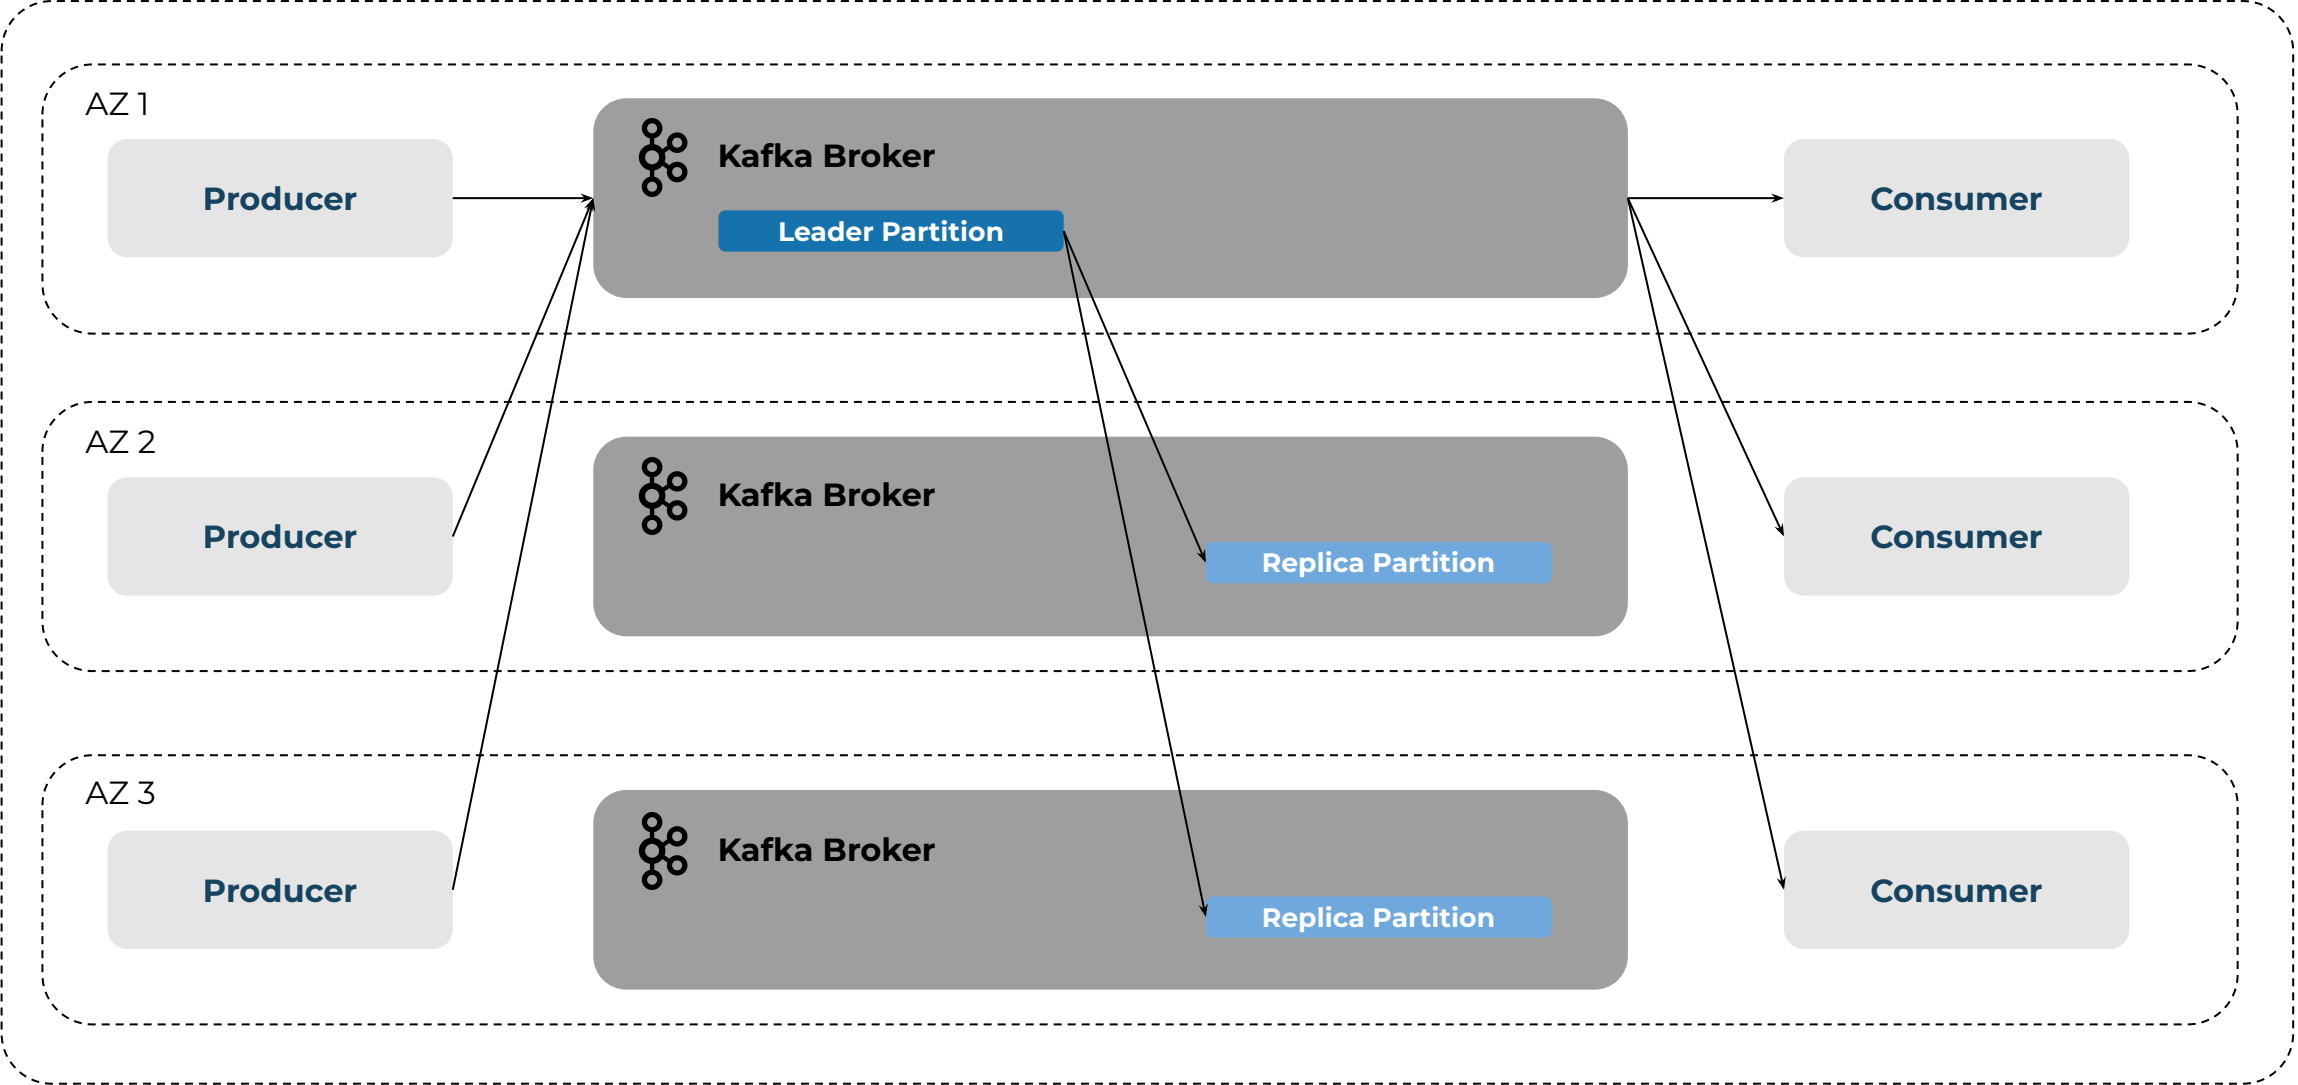


no tiered storage



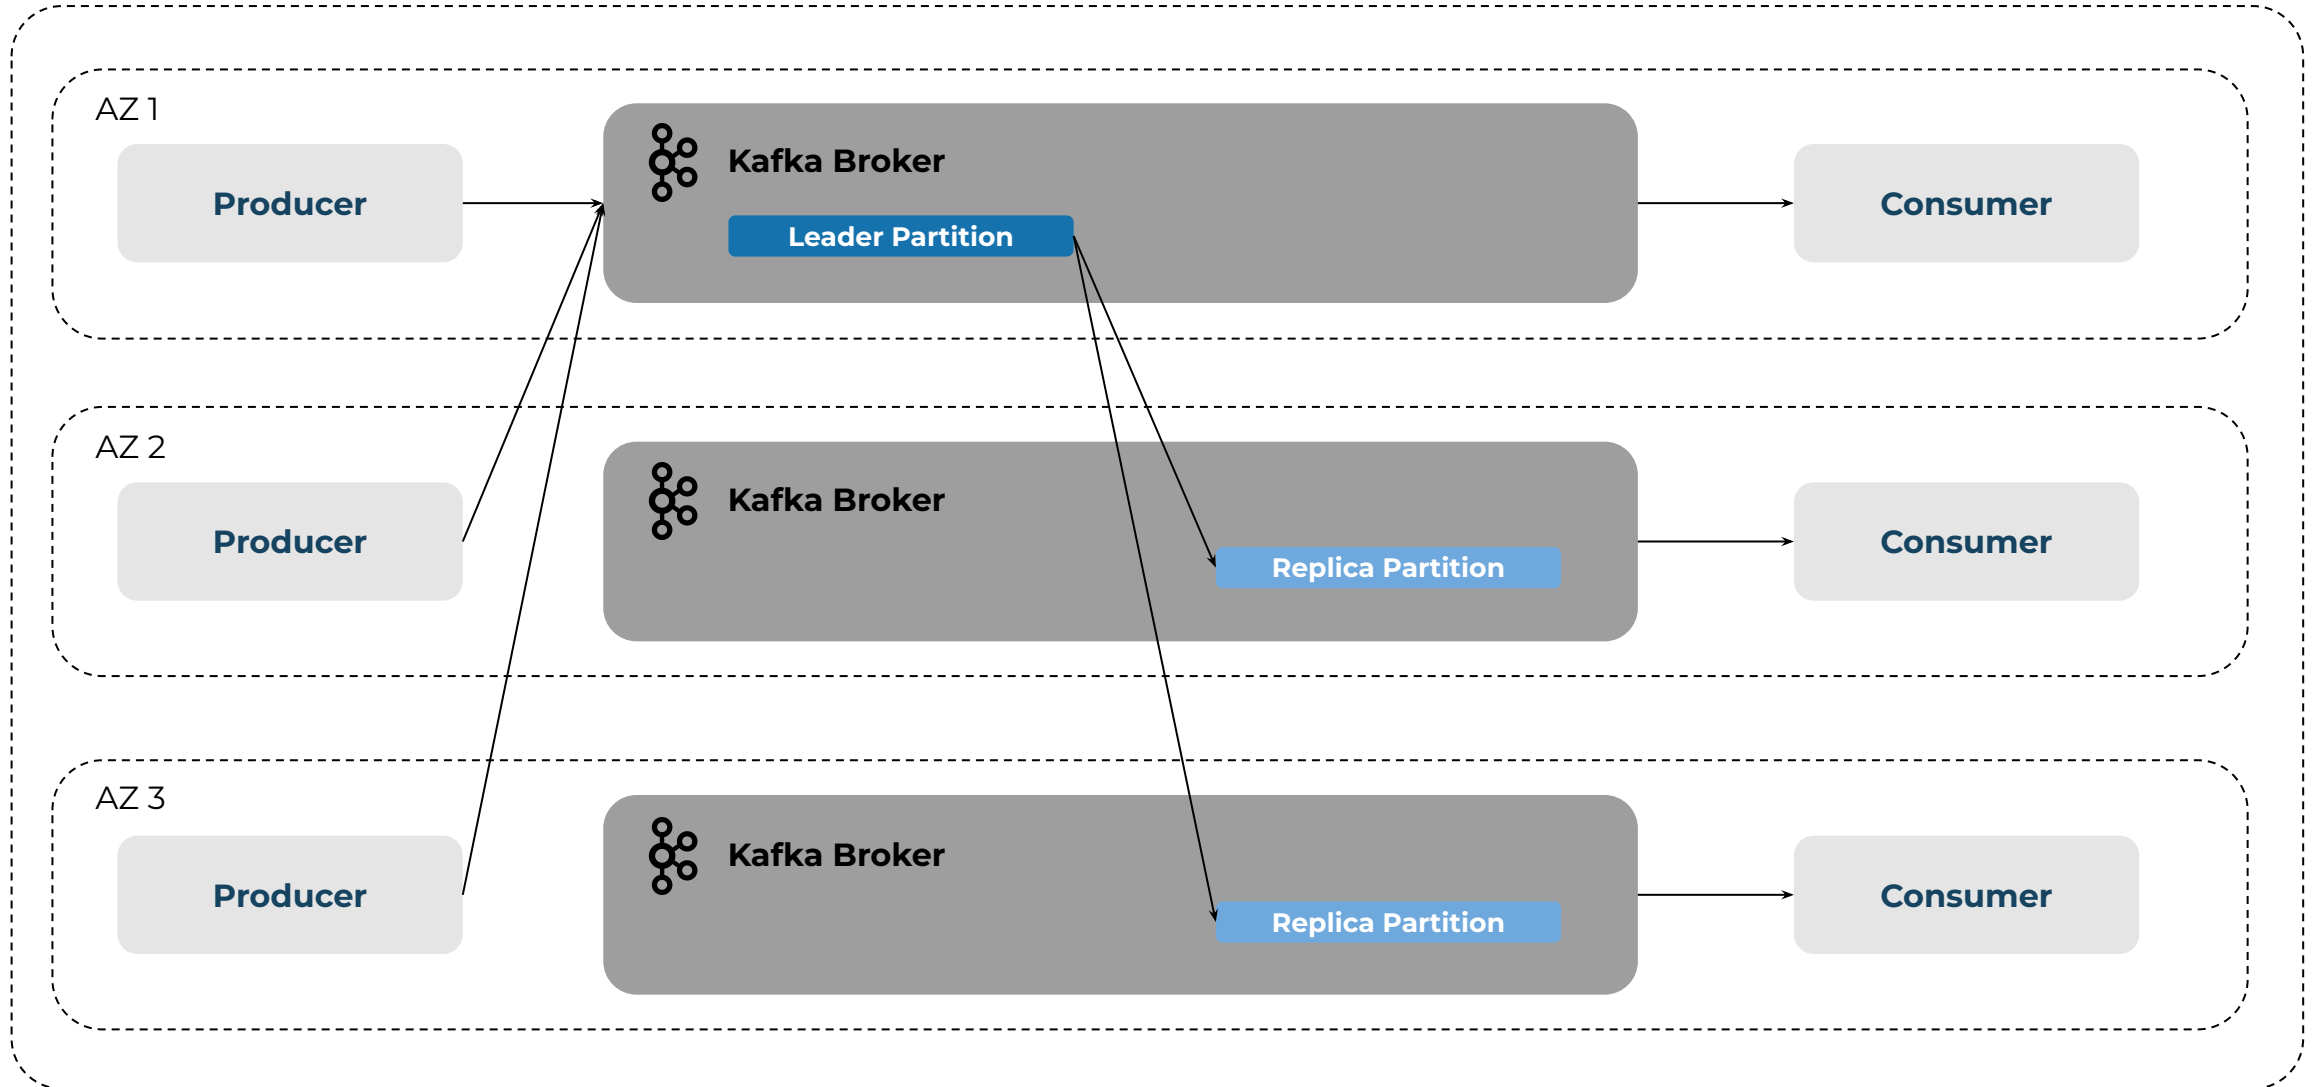
tiered storage

Cross-AZ networking in 2019



Cross-AZ networking in 2025

Fetch from follower and **rack-aware partition assignor** more or less solved the issue of consumer-side AZ traffic



Cross-AZ networking in 2025

Pricing differs among cloud vendors

0.02 \$/GB



0.01 \$/GB



0 \$/GB



Chapter 2:

Stateless Kafka brokers



What's “wrong” with kafka?



Apache Kafka
(w/o Tiered Storage)

- **Cross-az networking**
- **Expensive long-term storage**
- **Poor cluster elasticity**
- **High Operating complexity**



Apache Kafka
+ Tiered Storage

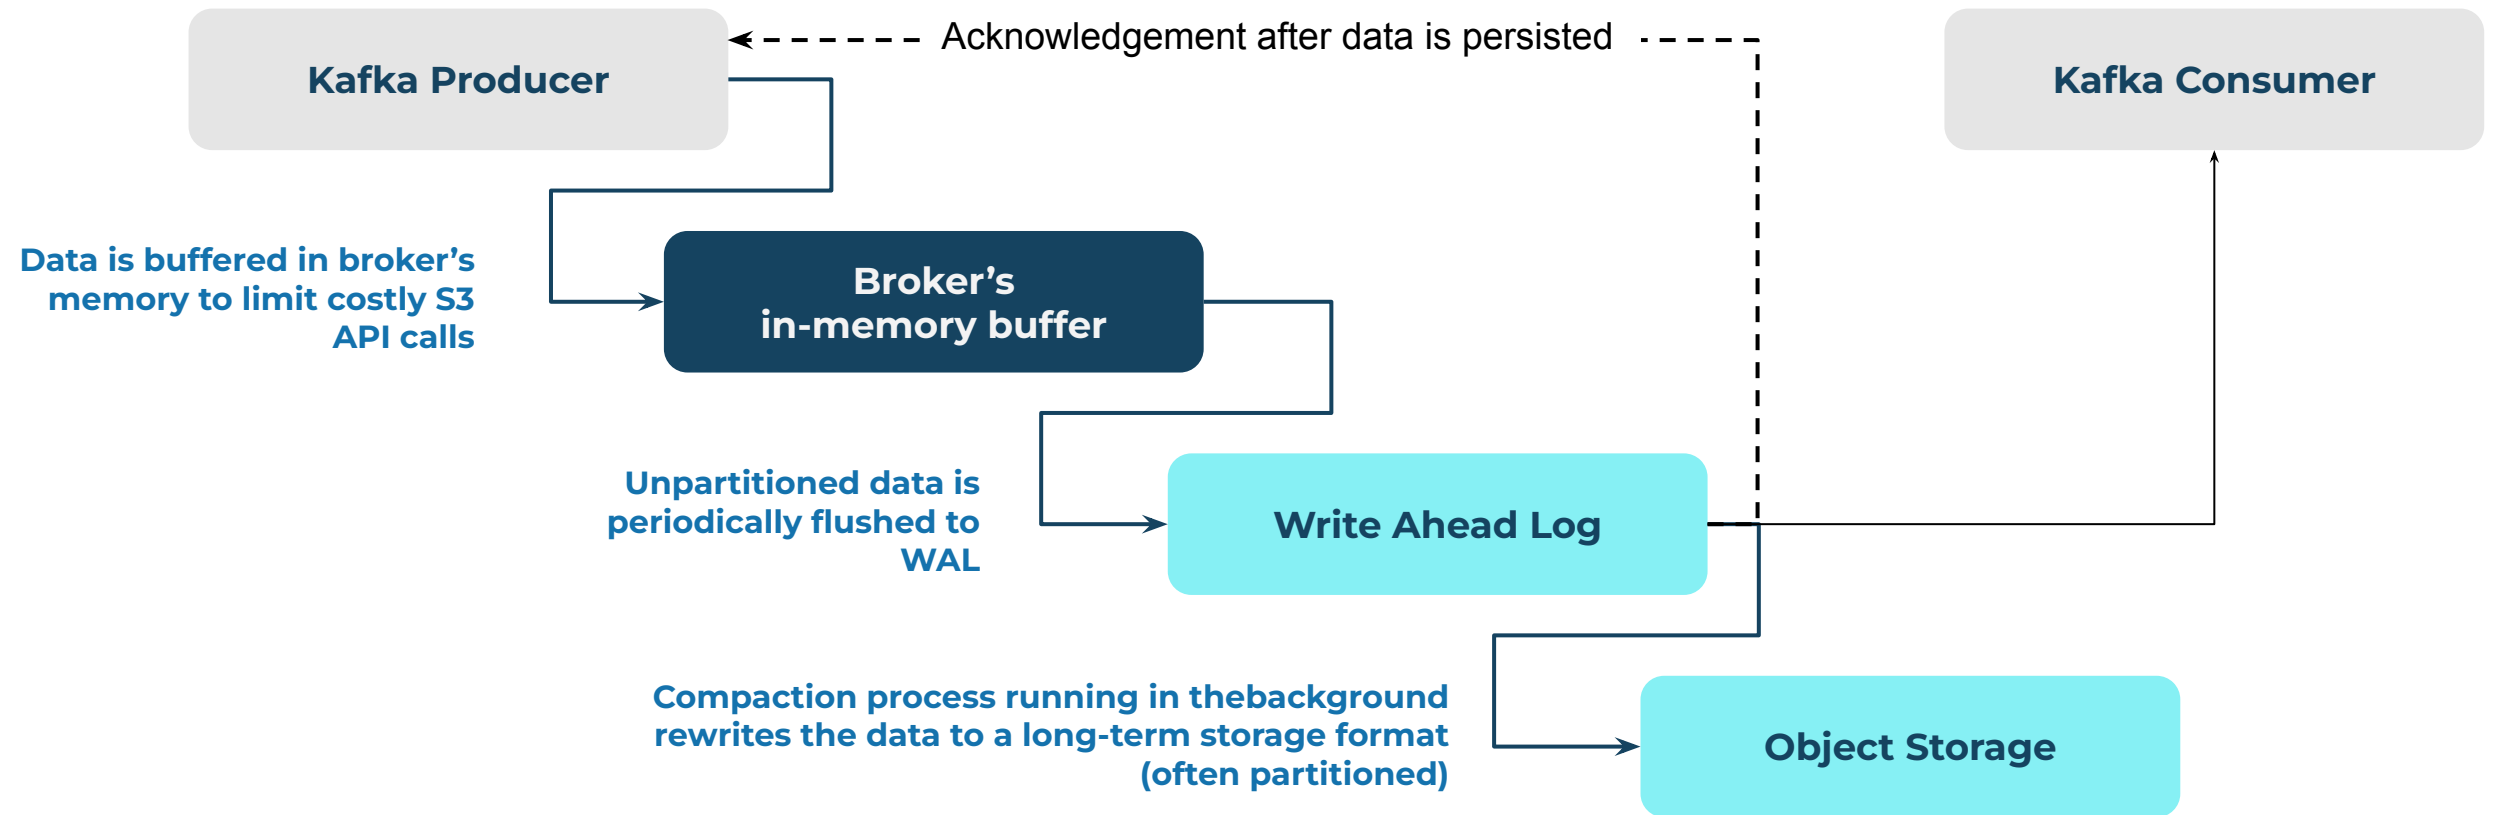
- **Cross-az networking**
- + **Cheap S3 storage**
- + **Elastic scaling**
- + **Simpler to operate**

Stateless Kafka
Brokers

- + **No cross-az networking**
(with exceptions)
- + **Cheap S3 storage**
- ++ **Extremely elastic scaling**
- ++ **Much simpler to operate**
- **Not battle-tested**
- **Higher latency**
(with exceptions)

What do they share?

Generic Stateless Kafka data flow



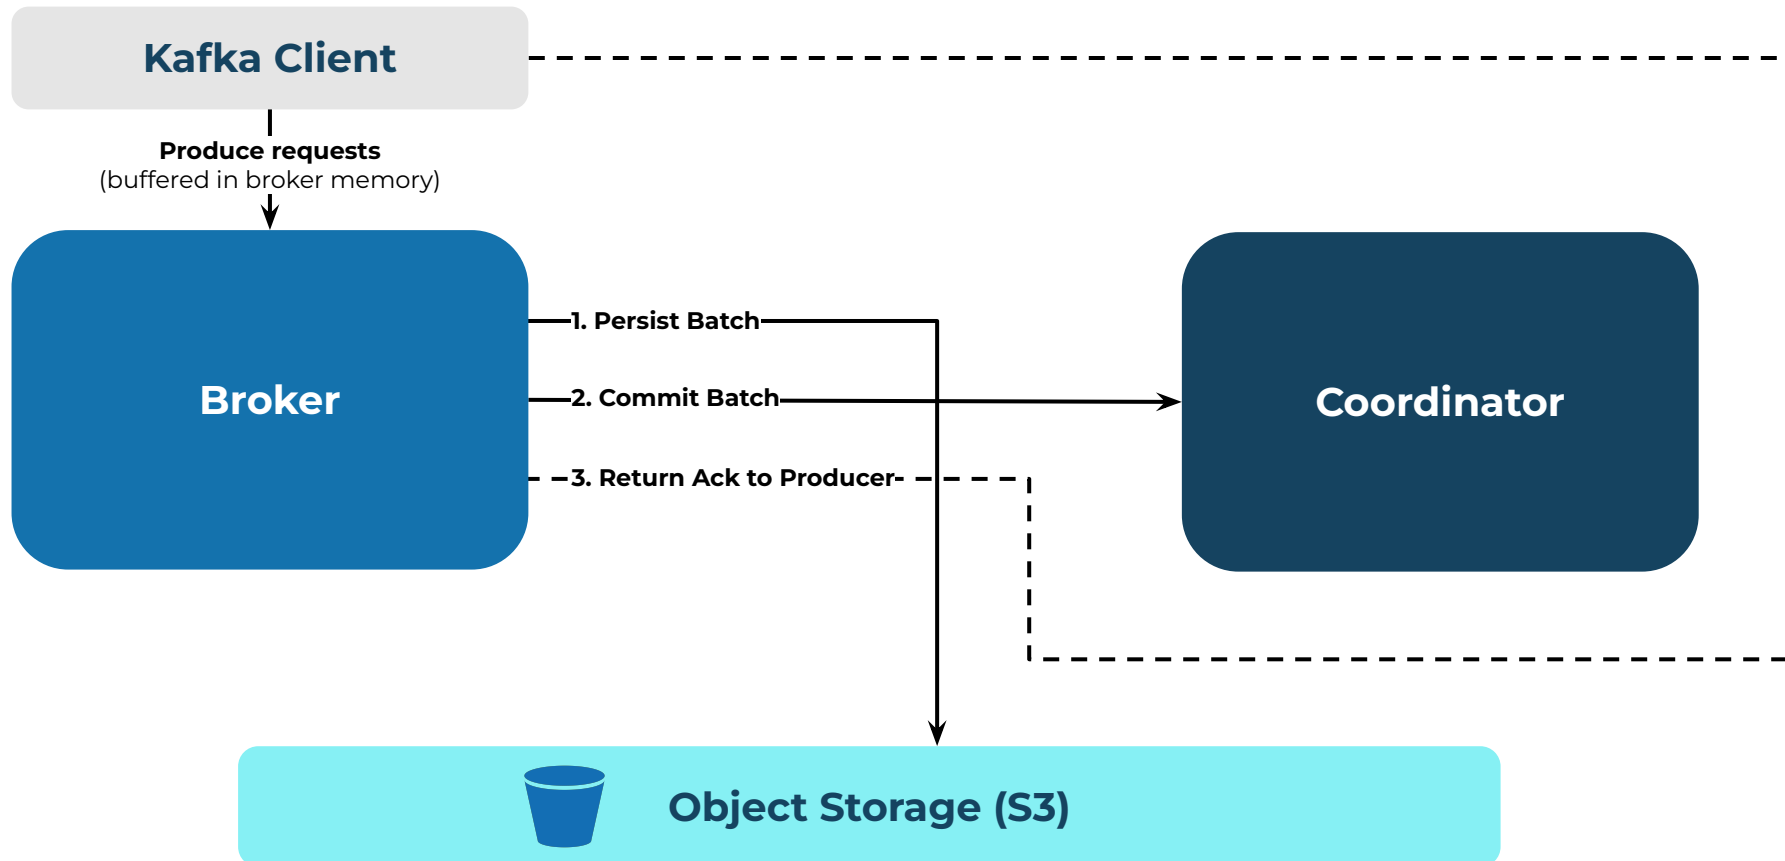
What do they share?

Popular features

- **Broker-side** schema validation
- Native **Iceberg** integration
- Schema registry included in binaries
- **Multi-region** active-active clusters
- Easy to migrate (easy to deploy drop-in replacements)


Coordinator based approach

- Leaderless - any client can read/write to any broker
- Batch contains records for all partitions
- Compaction process running in the background




Coordinator based approach

Who does it and how?




Diskless Kafka Topics in Apache Kafka?


KIP-1150 in a nutshell



Leaderless approach



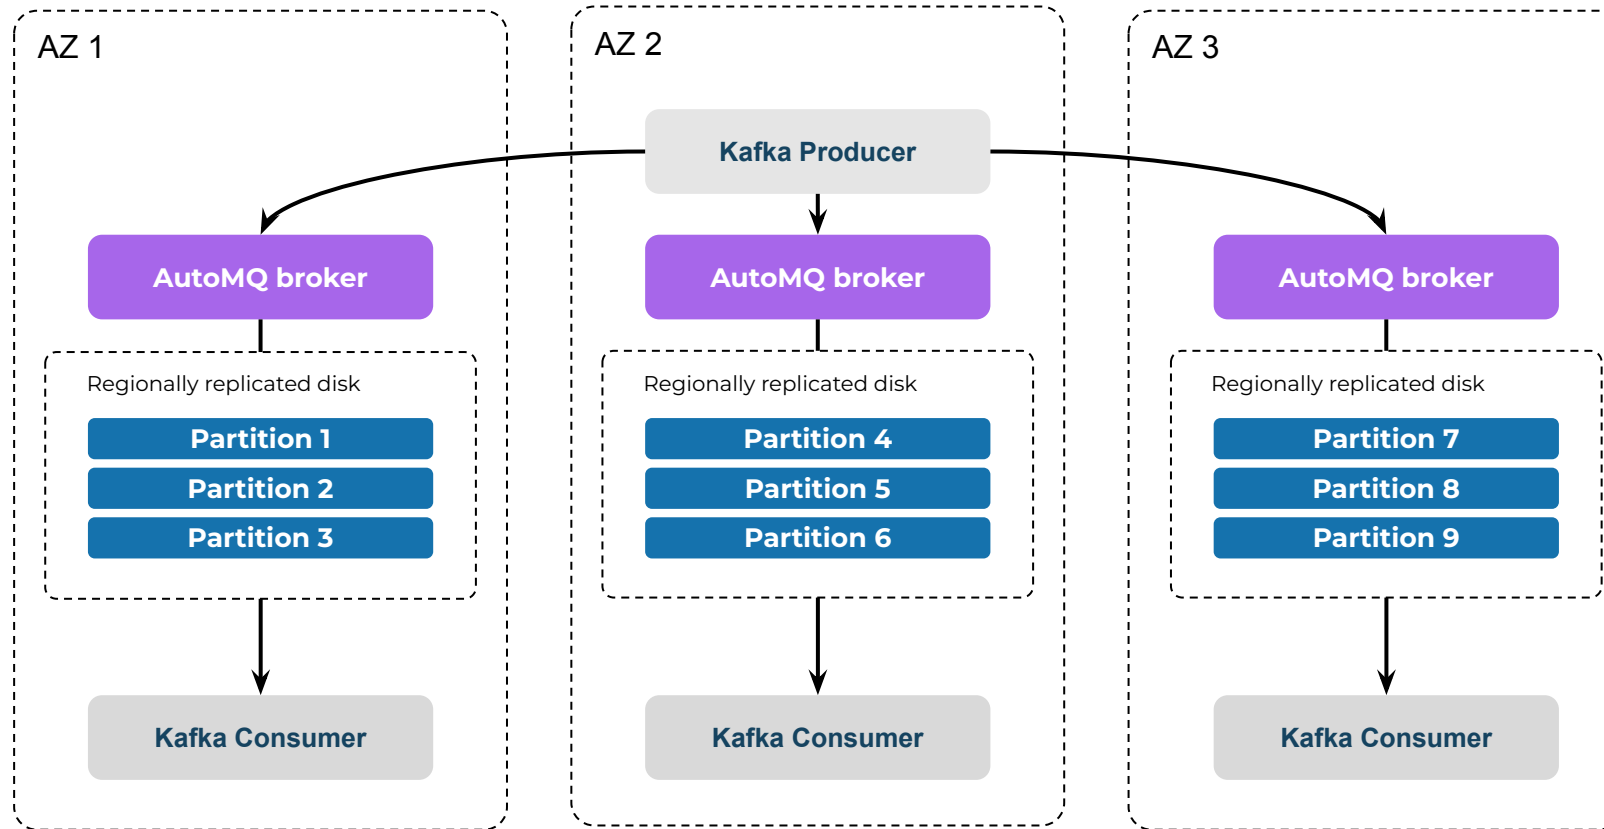
Batch Coordinator as a pluggable component with default implementation built as metadata topics using KRaft



KIP still under discussion

AutoMQ with Regional Block Storage

(example: Azure ZRS managed disk)



AutoMQ is a Kafka fork (Kafka 3.9 right now) with rewritten storage layer. Almost all of the code is open source, but this setup isn't.

AutoMQ Brokers are leaders for a partition with replication factor = 1

Data is synchronously replicated on the infrastructure level - in the case of AZ outage disk can be remounted in different AZ in ~30 seconds

Compaction process moves the data to S3 every few seconds

What's open source?



AutoMQ



Inkless



tansu

Chapter 3:

Cost and performance




Cost & latency comparison

Traffic I'm trying to normalize available data to:

- **1 GB/s write**
- **3 GB/s read**
- 4:1 compression ratio
- 7 days retention
- **Fetch from follower enabled**
- **AWS prices**

Sources:

- <https://www.redpanda.com/blog/kafka-kraft-vs-redpanda-performance-2023>
- <https://www.confluent.io/blog/kafka-vs-kora-latency-comparison/>
- <https://2minutestreaming.com/tools/apache-kafka-calculator/>
- <https://www.warpstream.com/pricing>
- <https://buf.build/docs/bufstream/cost/#optimized-for-cloud-economics>
- <https://www.automq.com/blog/automq-performance-benchmark-on-azure>
- <https://www.automq.com/pricing>





	Kafka w/o Tiered Storage	Kafka with Tiered Storage	Warpstream (II = Low Latency variant)	Bufstream	AutoMQ Enterprise	AutoMQ Open Source
EC2 Instances	41 970	11 658	4 811	4 448	3 510	2 470
Networking	145 590	145 596	0	904	35 040	52 (II: 35 040)
S3 API	-	-	4 620	4 508	1 101	6 670
Storage	52 623	17 931	13 900	13 900	13 900	13 900
Vendor fees	-	-	56 000	20 736	9 521	0
<hr/>						
Monthly cost	240 183	175 185	79 125 (~100k II)	42 625	61 201	23 092
Monthly cost w/o networking	94 593	29 589	79 125	41 721	26 161	23 039
<hr/>						
P50 E2E latency	19 ms	19 ms	500 ms (II: 180ms)	260 ms	9 ms	1.5 s II: 100-300ms
P99 E2E latency	43 ms	43 ms	900 ms (II: 600ms)	500 ms	40 ms	3.1 s

What's most valuable to us?



Drastic reduction of cross-AZ networking traffic

(very impactful unless you're on Azure)




Operational and architectural simplicity

Unlocking new markets for Kafka-compatible systems



Innovations in Data Streaming

Q & A?



Jan Siekierski

<https://www.linkedin.com/in/jan-siekierski/>

<https://kentra.io>

That's all, folks!